

BT-3102

COMPUTATIONAL METHODS FOR BUSINESS ANALYTICS

L12 Markov Decision Process (MDP) & Reinforcement Learning (RL)

Aditya Karanam

Optional Readings

- ▶ *N04-MDP-RL.pdf* (on Canvas)
- ▶ AIMA Chapters 17.1 to 17.3
 - ▶ *Sequential Decision Problems; Value Iteration; Policy Iteration*
- ▶ *Reinforcement Learning: An Introduction*, Richard S. Sutton and Andrew G. Barto, 2018. MIT Press. (Selected chapters freely available online)
- ▶ *Algorithms to Live By: The Computer Science of Human Decisions*, Brian Christian and Tom Griffiths, 2017. Picador

Algorithms to Live By

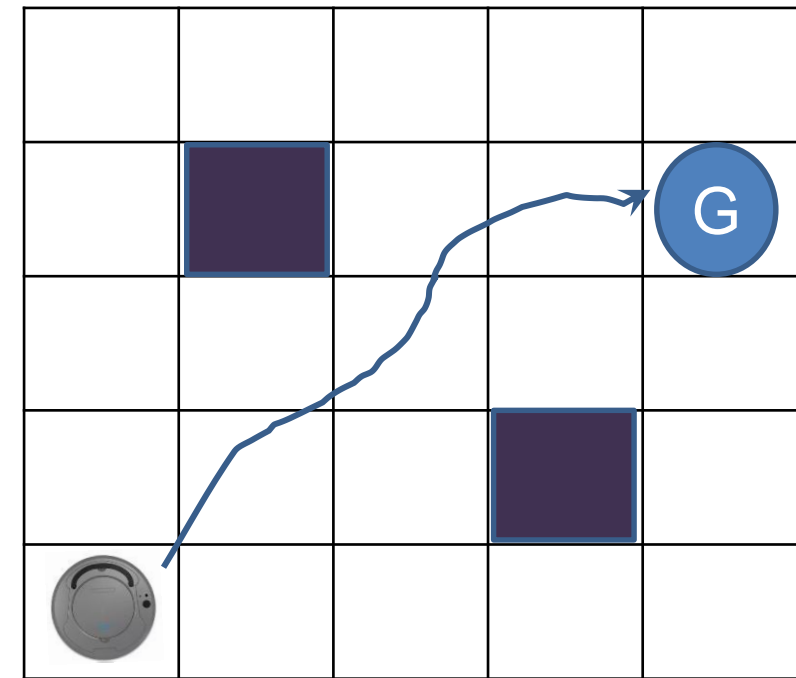


The
COMPUTER SCIENCE
of
HUMAN DECISIONS

Brian Christian and Tom Griffiths

Motivation

- ▶ Robot can move front, back, left, and right.
- ▶ It is now placed in a square space and our objective is to train it to reach a goal (or charging station) irrespective of its position
 - ▶ Assume that space is discrete, i.e., a grid
 - ▶ The state of the robot is its position in the grid
- ▶ The first step is to model the movement or transitions
 - ▶ How do we do that?




Transition Probability

- ▶ $T(s, a, s') = p(s'|a, s)$ gives the probability of going from s to s' upon taking action a
 - ▶ E.g., robot can select an action $a \in A = \{\text{front, back, left, right}\}$
- ▶ Assume that it follows first order Markov property
 - ▶ $p(s_t|a_{t-1}, s_{t-1}, a_{t-2}, s_{t-2}, \dots, a_0, s_0) = p(s_t|a_{t-1}, s_{t-1})$
- ▶ Action a is *non-deterministic*
 - ▶ E.g., when robot chooses to go front,
it moves front with probability 0.8, and
it moves left with probability 0.1, and
it moves right with probability 0.1

What else is required?

- ▶ Robot can reach its goal by taking:
 {Front, Front, Left, Left, Left, Left, Front}
- ▶ Consider the probability for each of these actions is 0.9
- ▶ Probability it reaches the goal: $0.9^7 = 0.48$
- ▶ Apart from the transition probability,
 - ▶ Also, need to communicate or incentivize the robot to achieve what we want



Reward and Utility

- ▶ *Reward* is a way of communicating to the robot (or agent) what you want it to achieve and ask it to maximize the *cumulative reward in the long run*.
 - ▶ $R(s, a, s')$ is a reward function that gives the reward at each state s
 - ▶ *Assumption*: Reward is obtained after taking action a and moving out of state s
- ▶ Examples:
 - ▶ To make a robot learn to walk, the reward for each time step can be proportional to the robot's forward motion
 - ▶ In making a robot learn to escape from a maze, the reward is often -1 for every time step it spends inside the maze
 - ▶ How would you design the reward to a robot in Kentridge Cafe so that it learns to deliver the food to the table?

Utility Function and Discounting Factor

- ▶ The *utility function* over a path $s_0, s_1, s_2 \dots$ gives the sum of discounted rewards,
- ▶ $U[s_0, s_1, s_2, \dots] = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t R(s_t)$
 - ▶ γ : discounting factor with value in $(0, 1)$
 - ▶ Rewards obtained later are worth much less today
- ▶ Utility of a path is bounded:
 - ▶ $U[s_0, s_1, s_2, \dots] \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = \frac{R_{max}}{1 - \gamma}$

Markov Decision Process (MDP)

- ▶ An MDP is defined by
 - ▶ A set of states S
 - ▶ A set of actions A
 - ▶ A transition probability function $T(s, a, s') = p(s'|a, s)$
 - ▶ A reward function $R(s, a, s')$
 - ▶ Discount: γ (*excluded sometimes*)

MDP: Policy Function

- ▶ Goal is to learn an *optimal* policy $\pi^*: S \rightarrow A$
 - ▶ Policy: A function π that specifies an action for each state
 - ▶ *Optimal* policy ($\pi^*(s)$): a function that specifies the action to take in state s that results in the highest *expected* utility
- ▶ Why *expected*?

Notations

- ▶ $V^*(s)$: value of state s * : represents the optimal value
 - ▶ i.e., the expected utility of starting in state s and acting optimally thereafter
- ▶ $Q^*(s, a)$: Q-value of state s and action a
 - ▶ i.e., the expected utility of starting in state s , taking action a and acting optimally thereafter
- ▶ $\pi^*(s)$: The optimal policy
 - ▶ i.e., it specifies the action that should be taken in state s in order to obtain the highest expected utility
 - ▶ Assumes that you act optimally at every step.

Relationship between V^* , Q^* , π^*

1. $V^*(s) = Q^*(s, a)$

2. $Q^*(s, a) =$

3.
$$\max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$
$$= \sum_{s'} T(s, \pi^*(s), s') [R(s, \pi^*(s), s') + \gamma V^*(s')]$$

How to express optimal policy $\pi^*(s)$?

4. $\pi^*(s) =$
 $=$

Value Iteration Algorithm

- ▶ $V_i^*(s)$ is the estimate of $V^*(s)$ in iteration i
- 1. Start with $V_0^*(s) = 0 \quad \forall s \in S$
- 2. Compute V_{i+1}^* given V_i^* :
$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')] \quad (\text{Bellman equation})$$
- 3. Repeat until convergence
 - ▶ i.e., $V_{i+1}^*(s)$ is very close to $V_i^*(s) \quad \forall s$

Convergence Proof of Value Iteration: Basic Idea

- ▶ Consider a simple MDP with a single state and a single action

- ▶ $V_{i+1}^* = R(s) + \gamma V_i^*$

- ▶ For true optimal V^* , $V^* = R(s) + \gamma V^*$

- ▶ Subtracting the above two equations:

$$V_{i+1}^* - V^* = \gamma(V_i^* - V^*) \quad (\gamma < 1)$$

- ▶ Distance between the estimated and optimal would decrease by a factor γ for each iteration

Practice - Value Iteration

- ▶ $S = \{0, 1, 2, 3\}$; $A = \{M, B\}$ (M : mini step; B : big step)

0	1	2	3
---	---	---	---

- ▶ Take mini step, move forward by 1 with prob 1.0
- ▶ Take big step, move forward by 2 with prob 0.3 and stay in same place with prob 0.7
- ▶ Can never exceed state 3. If you take a step and exceed state 3, then you will stay in state 3. If you take any action in state 3, you will always remain in state 3
- ▶ $T(k, M, k + 1) = \quad \forall k \in \{0, 1, 2\}$
 - ▶ $T(3, M, 3) =$
 - ▶ $T(k, B, k + 2) = \quad T(k, B, k) = \quad \forall k \in \{0, 1\}$
 - ▶ $T(2, B, 3) = \quad T(2, B, 2) = \quad T(3, B, 3) =$
- ▶ $R(s, a, s') = |s - s'|$
- ▶ What are the values of $V_1^*(s) \forall s$? (Let $\gamma = 0.5$)

Practice: Value of $V_1^*(0)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]]$
- ▶ $V_1^*(0) \leftarrow$
 $\max \{ T(0, M, 1)[R(0, M, 1) + \gamma V_0^*(1)],$
 $T(0, B, 2)[R(0, B, 2) + \gamma V_0^*(2)] + T(0, B, 0)[R(0, B, 0) + \gamma V_0^*(0)] \}$
- ▶ $V_1^*(0) \leftarrow \max \{ 1.0 \cdot [1 + 0.5 \cdot 0],$
 $0.3 \cdot [2 + 0.5 \cdot 0] + 0.7 \cdot [0 + 0.5 \cdot 0] \}$
- ▶ $V_1^*(0) \leftarrow 1.0$

Practice: Value of $V_1^*(1)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]]$
- ▶ $V_1^*(1) \leftarrow$
 $\max \{ T(1, M, 2) [R(1, M, 2) + \gamma V_0^*(2)],$
 $T(1, B, 3) [R(1, B, 3) + \gamma V_0^*(3)] + T(1, B, 1) [R(1, B, 1) + \gamma V_0^*(1)] \}$
- ▶ $V_1^*(1) \leftarrow \max \{ 1.0 \cdot [1 + 0.5 \cdot 0],$
 $0.3 \cdot [2 + 0.5 \cdot 0] + 0.7 \cdot [0 + 0.5 \cdot 0] \}$
- ▶ $V_1^*(1) \leftarrow 1.0$

Practice: Value of $V_1^*(2)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]]$
- ▶ $V_1^*(2) \leftarrow$
 $\max \{ T(2, M, 3) [R(2, M, 3) + \gamma V_0^*(3)],$
 $T(2, B, 3) [R(2, B, 3) + \gamma V_0^*(3)] + T(2, B, 2) [R(2, B, 2) + \gamma V_0^*(2)] \}$
- ▶ $V_1^*(2) \leftarrow \max \{ 1.0 \cdot [1 + 0.5 \cdot 0],$
 $0.3 \cdot [1 + 0.5 \cdot 0] + 0.7 \cdot [0 + 0.5 \cdot 0] \}$
- ▶ $V_1^*(2) \leftarrow 1.0$

Practice: Value of $V_1^*(3)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$
- ▶ $V_1^*(3) \leftarrow \max \left\{ \begin{array}{l} T(3, M, 3) [R(3, M, 3) + \gamma V_0^*(3)], \\ T(3, B, 3) [R(3, B, 3) + \gamma V_0^*(3)] \end{array} \right\}$
- ▶ $V_1^*(3) \leftarrow \max \left\{ \begin{array}{l} 1.0 \cdot [0 + 0.5 \cdot 0], \\ 1.0 \cdot [0 + 0.5 \cdot 0] \end{array} \right\}$
- ▶ $V_1^*(3) \leftarrow 0.0$
- ▶ $V_1^*(0) = 1.0, V_1^*(1) = 1.0, V_1^*(2) = 1.0, V_1^*(3) = 0.0$

Quiz: Value of $V_2^*(1)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$
- ▶ $V_2^*(1) = ?$

Practice: Value of $V_2^*(2)$

- ▶ $V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]]$
- ▶ $V_2^*(2) \leftarrow$
 $\max \{ T(2, M, 3) [R(2, M, 3) + \gamma V_1^*(3)],$
 $T(2, B, 3) [R(2, B, 3) + \gamma V_1^*(3)] + T(2, B, 2) [R(2, B, 2) + \gamma V_1^*(2)] \}$
- ▶ $V_2^*(2) \leftarrow \max \{ 1.0 \cdot [1 + 0.5 \cdot 0],$
 $0.3 \cdot [1 + 0.5 \cdot 0] + 0.7 \cdot [0 + 0.5 \cdot 1] \}$
- ▶ $V_2^*(2) \leftarrow 1.0$

Practice: Value of $V_2^*(3)$

$$\triangleright V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$$

$$\triangleright V_2^*(3) \leftarrow \max \left\{ \begin{array}{l} T(3, M, 3) [R(3, M, 3) + \gamma V_1^*(3)], \\ T(3, B, 3) [R(3, B, 3) + \gamma V_1^*(3)] \end{array} \right\}$$

$$\triangleright V_2^*(3) \leftarrow \max \left\{ \begin{array}{l} 1.0 \cdot [0 + 0.5 \cdot 0], \\ 1.0 \cdot [0 + 0.5 \cdot 0] \end{array} \right\}$$

$$\triangleright V_2^*(3) \leftarrow 0.0$$

$$V_2^*(0) = 1.5, \quad V_2^*(1) = 1.5, \quad V_2^*(2) = 1.0, \quad V_2^*(3) = 0.0$$

Practice: $\pi^*(s)$?

- ▶ Assume we obtained convergence: $V^* = V_2^*$
- ▶ $\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
- ▶ $\pi^*(0) = \operatorname{argmax}_a \sum_{s'} T(0, a, s') [R(0, a, s') + \gamma V^*(s')]$
 - ▶ $a = M$: $T(0, M, 1) [R(0, M, 1) + \gamma V^*(1)] = 1.0 [1 + 0.5 \cdot 1.5] = 1.75$
 - ▶ $a = B$: $T(0, B, 2) [R(0, B, 2) + \gamma V^*(2)] + T(0, B, 0) [R(0, B, 0) + \gamma V^*(0)]$
 $= 0.3 [2 + 0.5 \cdot 1.0] + 0.7 [0 + 0.5 \cdot 1.5] = 1.275$
- ▶ $\pi^*(1) = \operatorname{argmax}_a \sum_{s'} T(1, a, s') [R(1, a, s') + \gamma V^*(s')]$
 - ▶ $a = M$: $T(1, M, 2) [R(1, M, 2) + \gamma V^*(2)] = 1.0 [1 + 0.5 \cdot 1.0] = 1.5$
 - ▶ $a = B$: $T(1, B, 3) [R(1, B, 3) + \gamma V^*(3)] + T(1, B, 1) [R(1, B, 1) + \gamma V^*(1)]$
 $= 0.3 [2 + 0.5 \cdot 0] + 0.7 [0 + 0.5 \cdot 1.5] = 1.125$

How to find $Q^*(s, a)$?

► With Q-value iteration algorithm

1. Start with $Q_0^*(s, a) = 0, \forall s \in S, \forall a \in A$

2. Compute $Q_{i+1}^*(s, a)$ given $Q_i^*(s, a)$:

$$Q_{i+1}^*(s, a) \leftarrow \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_i^*(s', a')]$$

3. Repeat until convergence

► i.e., $Q_{i+1}^*(s, a)$ is very close to $Q_i^*(s, a) \quad \forall s \in S$

MDP Application: Health Care

- ▶ Optimal timing of living-donor liver transplant
 - ▶ States: patient health status
 - ▶ Actions: Transplant, Wait
 - ▶ Transition function: Obtained by the progression of the disease
 - ▶ Reward: total discounted life expectancy
 - ▶ Discount factor: 0.99
- ▶ <https://pubsonline.informs.org/doi/10.1287/mnsc.1040.0287>

MDP Application: Finance

- ▶ Consumption-Investment Problem
 - ▶ **States:** Wealth, i.e., amount of money owned
 - ▶ **Actions:** Amount of money invested in stocks, bonds (remainder is consumed)
 - ▶ **Reward function:** Amount of money consumed (or utility function based on money owned)
 - ▶ **Transition probability:** estimated historical performance of stocks, bonds
 - ▶ <https://www.minet.uni-jena.de/Marie-Curie-ITN/SMIF/talks/Baeuerle.pdf>

Reinforcement Learning

- ▶ Given states S and actions A
- ▶ But don't know transition probs $T(s, a, s')$
- ▶ When agent (e.g., robot) in state s , it can experience a reward of getting to that state
- ▶ How to find $Q^*(s, a)$?

Naïve Sampling Approach to Estimate $Q^*(s, a)$

- ▶ Recall: $Q_{i+1}^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_i^*(s', a')]$
- ▶ Estimate $Q_{i+1}^*(s, a)$ using $Q_i^*(s, a)$ and samples obtained by:
 - ▶ Repeatedly taking one action at a time, and observing the reward and next state
 - ▶ Pretend can restart at s after action a
- ▶ Sample 1: $R(s, a, s'_1) + \gamma \max_{a'} Q_i^*(s'_1, a')$
- ▶ Sample 2: $R(s, a, s'_2) + \gamma \max_{a'} Q_i^*(s'_2, a')$
- ▶ ...
- ▶ Sample k : $R(s, a, s'_k) + \gamma \max_{a'} Q_i^*(s'_k, a')$
- ▶ $Q_{i+1}^*(s, a) = \frac{1}{k} \sum_{j=1}^k R(s, a, s'_j) + \gamma \max_{a'} Q_i^*(s'_j, a')$

Naïve Sampling: Analysis

- ▶ Problem with naïve sampling:
 - ▶ Extremely slow
 - ▶ Q-value is updated after every k samples or after k moves
- ▶ In practice, we collect the samples when we actually move.
 - ▶ Update Q-value after every move (s, a)
 - ▶ After every move, the agent (e.g., robot) observes the next state s' and experiences reward, $R(s, a, s')$

Q-Learning Algorithm

Q-Learning Algorithm:

1. For $i = 1, 2, 3 \dots$ (till convergence)
2. Collect a sample: s, a, s' and $R(s, a, s')$
3. Update running average of Q-values

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha[R(s, a, s') + \gamma \max_{a'} Q(s', a')]$$

- ▶ α : learning rate
 - ▶ Usually, start with $\alpha = 1$, and slowly shrink it to 0 as the number of iterations increases
- ▶ How to choose actions or collect samples?

How to choose action a in Q-learning?

- ▶ With probability ϵ , pick a randomly (exploration)
- ▶ With probability $1 - \epsilon$, pick a that maximizes current estimate of $Q(s, a)$,
i.e., $\operatorname{argmax}_a Q(s, a)$ (exploitation)
- ▶ Decrease ϵ with time as we gather more samples and obtain a better estimate of $Q(s, a)$

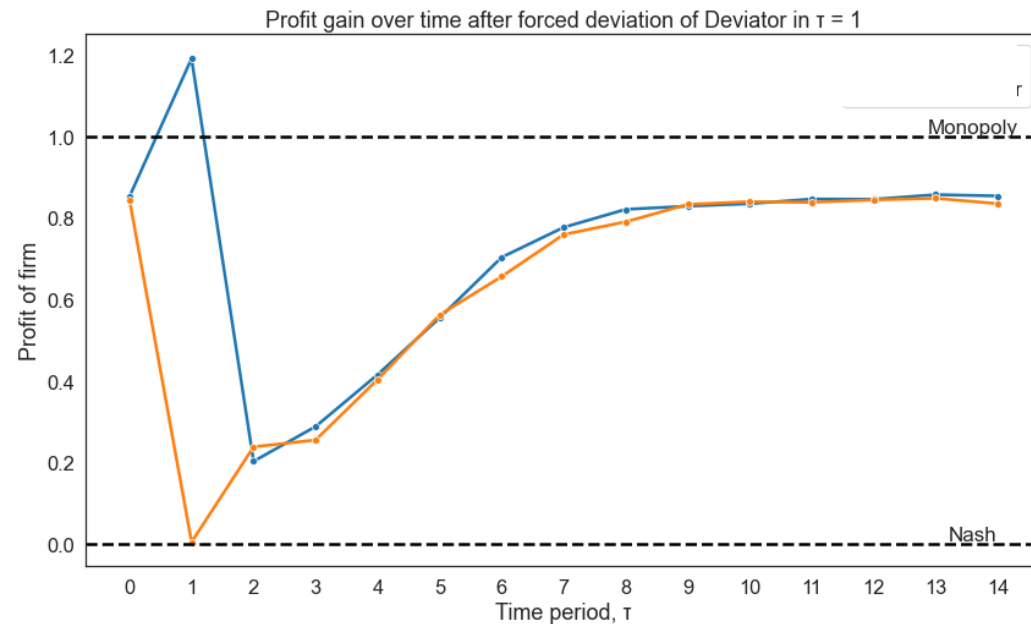
Applications

- ▶ Alpha Go: Deep Mind
 - ▶ Documentary: <https://www.youtube.com/watch?v=WXuK6gekU1Y>
- ▶ ChatGPT updates its responses
 - ▶ <https://huggingface.co/blog/rlhf>
- ▶ Trading, Advertising, Healthcare ...


Applications: RL algorithms for pricing

- Solve these questions by examining the behavior of Q-learning and Deep Q-Learning algorithms in repeated repeated economic games (pricing and auctions)

Preliminary results (work in progress)



E-commerce



Auctions
(First and Second Price Auctions)

Thank You
